



BREAST CANCER CLASSIFICATION

GROUP 7:

ASTHA TIBREWAL

HARSHIT GUPTA

PAYAL DANGI

KAUSTAV DASGUPTA

OBJECTIVE

We aimed at building a breast cancer classification model that could accurately and efficiently detect breast cancer. A bigger motivation was to see how healthcare professionals make informed decisions for diagnosis and treatment, leading to improved patient outcomes and reduced mortality rates.

The goal was to create a model that can accurately distinguish between breast tumors based on input features extracted from medical imaging data, such as ultrasound scans.

DATASET USED

<http://archive.ics.uci.edu/ml/datasets/breast+cancer+wisconsin+%28diagnostic%29>

This has 569 training examples on which we trained our model.

OUR METHOD

First of all we imported the dataset from the SKlearn databases and converted the data into a dataframe. After that we cleaned the data by checking for any NULL entries and scaling the values using MinMaxScaler. Lastly we splitted the data into training and testing data and trained a logistic regression model based on the training data and tested the trained model on the testing data.

Finally we got an accuracy of 98% which shows that the model was quite accurate.

TECHNIQUES IMPLEMENTED

Firstly we used KNN model to test the data, the KNN model gave us an accuracy of 96%. So in order to increase the accuracy we tried to use different models.

SVM model gave an accuracy of 97%

And the logistic regression model gave an overall accuracy of 98%, as the logistic model had the greatest accuracy we went with the logistic regression model.

CODE

In [25]:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.metrics import accuracy_score
from sklearn.metrics import roc_curve, auc
from sklearn import metrics
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import cohen_kappa_score
from sklearn.metrics import matthews_corrcoef
import scikitplot as skplt
```

In [26]:

```
from sklearn.datasets import load_breast_cancer
cancer = load_breast_cancer() # embeded dataset
df = pd.DataFrame(np.c_[cancer['target'], cancer['data']],
                  columns=np.append(['MB'], cancer['feature_names']))
df
```

Out[26]:

	MB	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	...	worst radius	worst texture	worst perimeter	worst area	worst smoothness	v
0	0.0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.30010	0.14710	0.2419	...	25.380	17.33	184.60	2019.0	0.16220	0.6
1	0.0	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.08690	0.07017	0.1812	...	24.990	23.41	158.80	1956.0	0.12380	0.1
2	0.0	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.19740	0.12790	0.2069	...	23.570	25.53	152.50	1709.0	0.14440	0.4
3	0.0	11.42	20.38	77.58	386.1	0.14250	0.28390	0.24140	0.10520	0.2597	...	14.910	26.50	98.87	567.7	0.20980	0.8
4	0.0	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.19800	0.10430	0.1809	...	22.540	16.67	152.20	1575.0	0.13740	0.2
...
564	0.0	21.56	22.39	142.00	1479.0	0.11100	0.11590	0.24390	0.13890	0.1726	...	25.450	26.40	166.10	2027.0	0.14100	0.2
565	0.0	20.13	28.25	131.20	1261.0	0.09780	0.10340	0.14400	0.09791	0.1752	...	23.690	38.25	155.00	1731.0	0.11660	0.1
566	0.0	16.60	28.08	108.30	858.1	0.08455	0.10230	0.09251	0.05302	0.1590	...	18.980	34.12	126.70	1124.0	0.11390	0.3

In [27]:

```
df.info()
# Checking for null entries if present

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   MB                                     569 non-null    float64
1   mean radius                           569 non-null    float64
2   mean texture                           569 non-null    float64
3   mean perimeter                         569 non-null    float64
4   mean area                             569 non-null    float64
5   mean smoothness                       569 non-null    float64
6   mean compactness                      569 non-null    float64
7   mean concavity                         569 non-null    float64
8   mean concave points                   569 non-null    float64
9   mean symmetry                         569 non-null    float64
10  mean fractal dimension                 569 non-null    float64
11  radius error                           569 non-null    float64
12  texture error                          569 non-null    float64
13  perimeter error                       569 non-null    float64
14  area error                            569 non-null    float64
15  smoothness error                      569 non-null    float64
16  compactness error                     569 non-null    float64
17  concavity error                       569 non-null    float64
18  concave points error                  569 non-null    float64
19  symmetry error                        569 non-null    float64
20  fractal dimension error                569 non-null    float64
21  worst radius                           569 non-null    float64
22  worst texture                          569 non-null    float64
23  worst perimeter                       569 non-null    float64
24  worst area                             569 non-null    float64
25  worst smoothness                      569 non-null    float64
26  worst compactness                     569 non-null    float64
27  worst concavity                       569 non-null    float64
28  worst concave points                   569 non-null    float64
29  worst symmetry                         569 non-null    float64
30  worst fractal dimension                569 non-null    float64
dtypes: float64(31)
memory usage: 137.9 KB
```

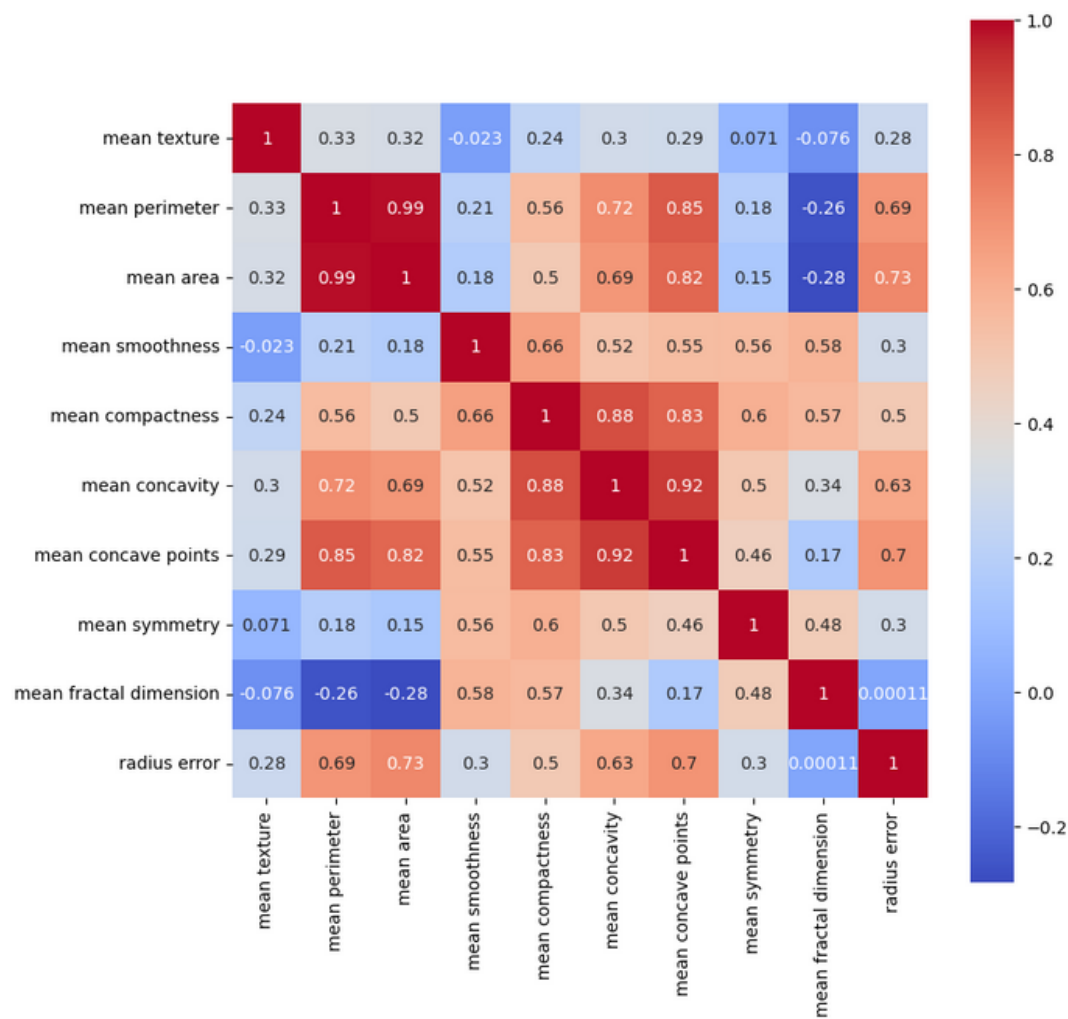
In [28]: `df.describe()`

Out[28]:

	MB	mean radius	mean texture	mean perimeter	mean area	mean smoothness	mean compactness	mean concavity	mean concave points	mean symmetry	...	worst radius	worst texture	perim
count	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	569.000000	...	569.000000	569.000000	569.000000
mean	0.627417	14.127292	19.289649	91.969033	654.889104	0.096360	0.104341	0.088799	0.048919	0.181162	...	16.269190	25.677223	107.269190
std	0.483918	3.524049	4.301036	24.298981	351.914129	0.014064	0.052813	0.079720	0.038803	0.027414	...	4.833242	6.146258	33.614129
min	0.000000	6.981000	9.710000	43.790000	143.500000	0.052630	0.019380	0.000000	0.000000	0.106000	...	7.930000	12.020000	50.400000
25%	0.000000	11.700000	16.170000	75.170000	420.300000	0.086370	0.064920	0.029560	0.020310	0.161900	...	13.010000	21.080000	84.170000
50%	1.000000	13.370000	18.840000	86.240000	551.100000	0.095870	0.092630	0.061540	0.033500	0.179200	...	14.970000	25.410000	97.610000
75%	1.000000	15.780000	21.800000	104.100000	782.700000	0.105300	0.130400	0.130700	0.074000	0.195700	...	18.790000	29.720000	125.400000
max	1.000000	28.110000	39.280000	188.500000	2501.000000	0.163400	0.345400	0.426800	0.201200	0.304000	...	36.040000	49.540000	251.200000

8 rows × 31 columns





```
In [30]: X = df.drop('MB', axis = 1)
y = df['MB']

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
In [31]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [32]: model_accuracy = {}
```

```
In [33]: from sklearn.linear_model import LogisticRegression
```

```
In [34]: a = list(range(-5, 10))
complexity_values = [10**i for i in a]
#xticks = list(range(1, len(complexity_values)+1))
param_values = {'C':complexity_values, 'penalty': ['l1', 'l2', 'elasticnet', 'none'], 'solver':['liblinear', 'newton-cg']}

clf = LogisticRegression()
m_lr = GridSearchCV(clf, param_grid=param_values, cv = 3, scoring = 'accuracy', return_train_score=True, n_jobs=-1)
m_lr.fit(X_train, y_train)

y_pred = m_lr.predict(X_test)

model_accuracy['Logistic Regression'] = accuracy_score(y_test, y_pred)

print('\n')
print('Prediction Accuracy: ', accuracy_score(y_test, y_pred))
print('\n')
print('confusion matrix: ')
print(confusion_matrix(y_test, y_pred))
print('\n')
print('classification report: ')
print(classification_report(y_test, y_pred))
```

```
fpr, tpr, threshold = roc_curve(y_test, y_pred)
auc = metrics.auc(fpr, tpr)
print("The AUC stats is: ", auc)
print("The kappa stats is: ", cohen_kappa_score(y_test, y_pred))
print("The MCC stats is: ", matthews_corrcoef(y_test, y_pred))
```

ROC curve

```
predicted_probab_lr = m_lr.predict_proba(X_test)
skplt.metrics.plot_roc(y_test, predicted_probab_lr)
```

Lift curve

```
skplt.metrics.plot_lift_curve(y_test, predicted_probab_lr)
plt.show()
```

```
best score: 0.9647960962007668
best parameters: {'C': 1000, 'penalty': 'l2', 'solver': 'newton-cg'}
best estimator: LogisticRegression(C=1000, solver='newton-cg')
```

Prediction Accuracy: 0.9824561403508771

confusion matrix:

```
[[42  1]
 [ 1 70]]
```

classification report:

	precision	recall	f1-score	support
0.0	0.98	0.98	0.98	43
1.0	0.99	0.99	0.99	71
accuracy			0.98	114
macro avg	0.98	0.98	0.98	114
weighted avg	0.98	0.98	0.98	114

The AUC stats is: 0.9813298395021289


The kappa stats is: 0.9626596790042581


The MCC stats is: 0.9626596790042581

The AUC stats is: 0.9813298395021289

The kappa stats is: 0.9626596790042581

The MCC stats is: 0.9626596790042581





```
5]: model_accuracy
```

```
5]: {'Logistic Regression': 0.9824561403508771}
```

RESULTS

We used regression to make the model and obtained an accuracy of 98%. The model was able to detect the type of breast cancer with 98% precision.

We also implemented the model using KNN and SVM. The accuracy we obtained for the KNN model was 96%.

CONCLUSION

This model aimed to develop an accurate and efficient system for distinguishing between benign and malignant breast tumors using machine learning techniques.

Different machine learning algorithms, such as logistic regression, support vector machines and KNNs were implemented and analyzed to give the best possible result. Our model was made by using regression techniques and further optimized to achieve the best possible performance.